

Modeling With and Without Outliers

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set_context('poster')
import random
```

```
In [2]: from project_helper import *
```

Import, Split and Standardize Data

```
In [3]: df = pd.DataFrame.from_csv('../data/merged/all_data_2006_to_2016.csv', index_col=None)
```

```
In [4]: # split train/test by indices before removing outliers
# so that we can have a consistent split between tests

# seed for consistent splits
random.seed(1636)

# all idxs
idxs = list(range(len(df)))

# idxs of outliers
outlier_idx = list((df.MSA_abbr == 'NEW_ORLEANS_LA').nonzero()[0])
outlier_idx += list(((df.MSA_abbr == 'MEMPHIS_TN') & (df.year == 2016)).nonzero()[0])
outlier_idx += list(((df.MSA_abbr == 'BATON_ROUGE_LA') & (df.year == 2007)).nonzero()[0])

# shuffle and split
random.shuffle(idxs)
split = int(len(idxs)*0.7)
train_all = idxs[:split]
test_all = idxs[split:]

# remove outliers
train_no = [idx for idx in train_all if idx not in outlier_idx]
test_no = [idx for idx in test_all if idx not in outlier_idx]

In [5]: # drop MSA names
df = df.drop(['MSA_corr', 'MSA_orig', 'MSA_abbr'], axis=1)
```

In [6]: *# separate labels and features*

```
label_col = 'murder_per_100_k'

# with outliers
x_train_all = df.iloc[train_all].drop([label_col], axis=1)
y_train_all = df.iloc[train_all][label_col]
x_test_all = df.iloc[test_all].drop([label_col], axis=1)
y_test_all = df.iloc[test_all][label_col]

# without outliers
x_train_no = df.iloc[train_no].drop([label_col], axis=1)
y_train_no = df.iloc[train_no][label_col]
x_test_no = df.iloc[test_no].drop([label_col], axis=1)
y_test_no = df.iloc[test_no][label_col]

print('All Sizes match: {}'.format(len(x_train_all)==len(y_train_all)))
print('Clean Sizes match: {}'.format(len(x_train_no)==len(y_train_no)))
```

All Sizes match: True
Clean Sizes match: True

In [7]: *# standardize data*

```
from sklearn.preprocessing import StandardScaler

# with outliers
standardizer = StandardScaler().fit(x_train_all)
x_train_all = standardizer.transform(x_train_all)
x_test_all = standardizer.transform(x_test_all)

# no outliers
standardizer = StandardScaler().fit(x_train_no)
x_train_no = standardizer.transform(x_train_no)
x_test_no = standardizer.transform(x_test_no)
```

```
In [8]: from sklearn.linear_model import LinearRegression, RidgeCV, LassoCV, BayesianRidge, HuberRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.svm import SVR
```

```
In [9]: # instantiate and fit models
def make_models(x_train, y_train):
    md = dict()

    md['linear'] = LinearRegression().fit(x_train, y_train)
    md['ridge'] = RidgeCV(cv=5).fit(x_train, y_train)
    md['huber'] = GridSearchCV(HuberRegressor(), {'epsilon': [1.0, 1.2, 1.4, 1.6, 1.8]}).fit(x_train, y_train)
    md['knn'] = GridSearchCV(KNeighborsRegressor(), {'n_neighbors': [5, 10, 20, 40]}).fit(x_train, y_train)
    md['adaboost'] = GridSearchCV(AdaBoostRegressor(), {'learning_rate': [0.1, 0.3, 0.6, 1.0]}).fit(x_train, y_train)
    md['svr'] = GridSearchCV(SVR(), {'C': [0.01, 0.1, 1, 10, 100], 'epsilon': [0.001, 0.01, 0.1, 1, 10]}).fit(x_train, y_train)

    return md
```

Fit Models

```
In [10]: # with outliers
exp_1 = run_experiment(make_models, 100, x_train_all, y_train_all, x_test_all, y_test_all)

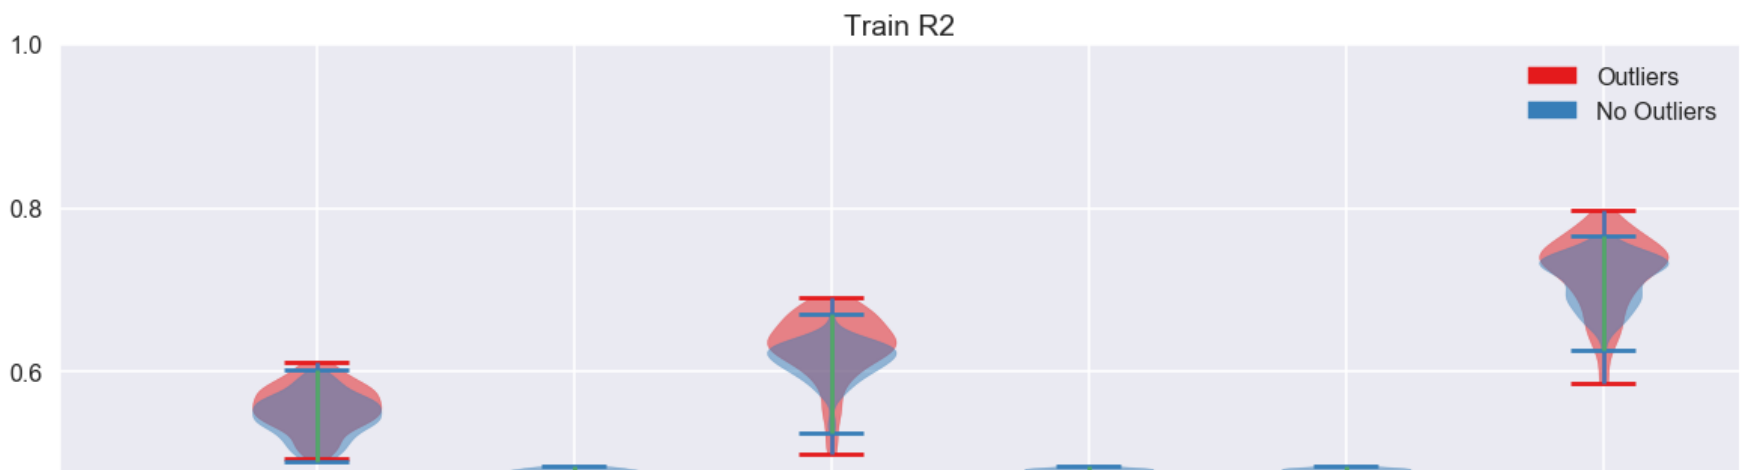
# without outliers
exp_2 = run_experiment(make_models, 100, x_train_no, y_train_no, x_test_no, y_test_no)

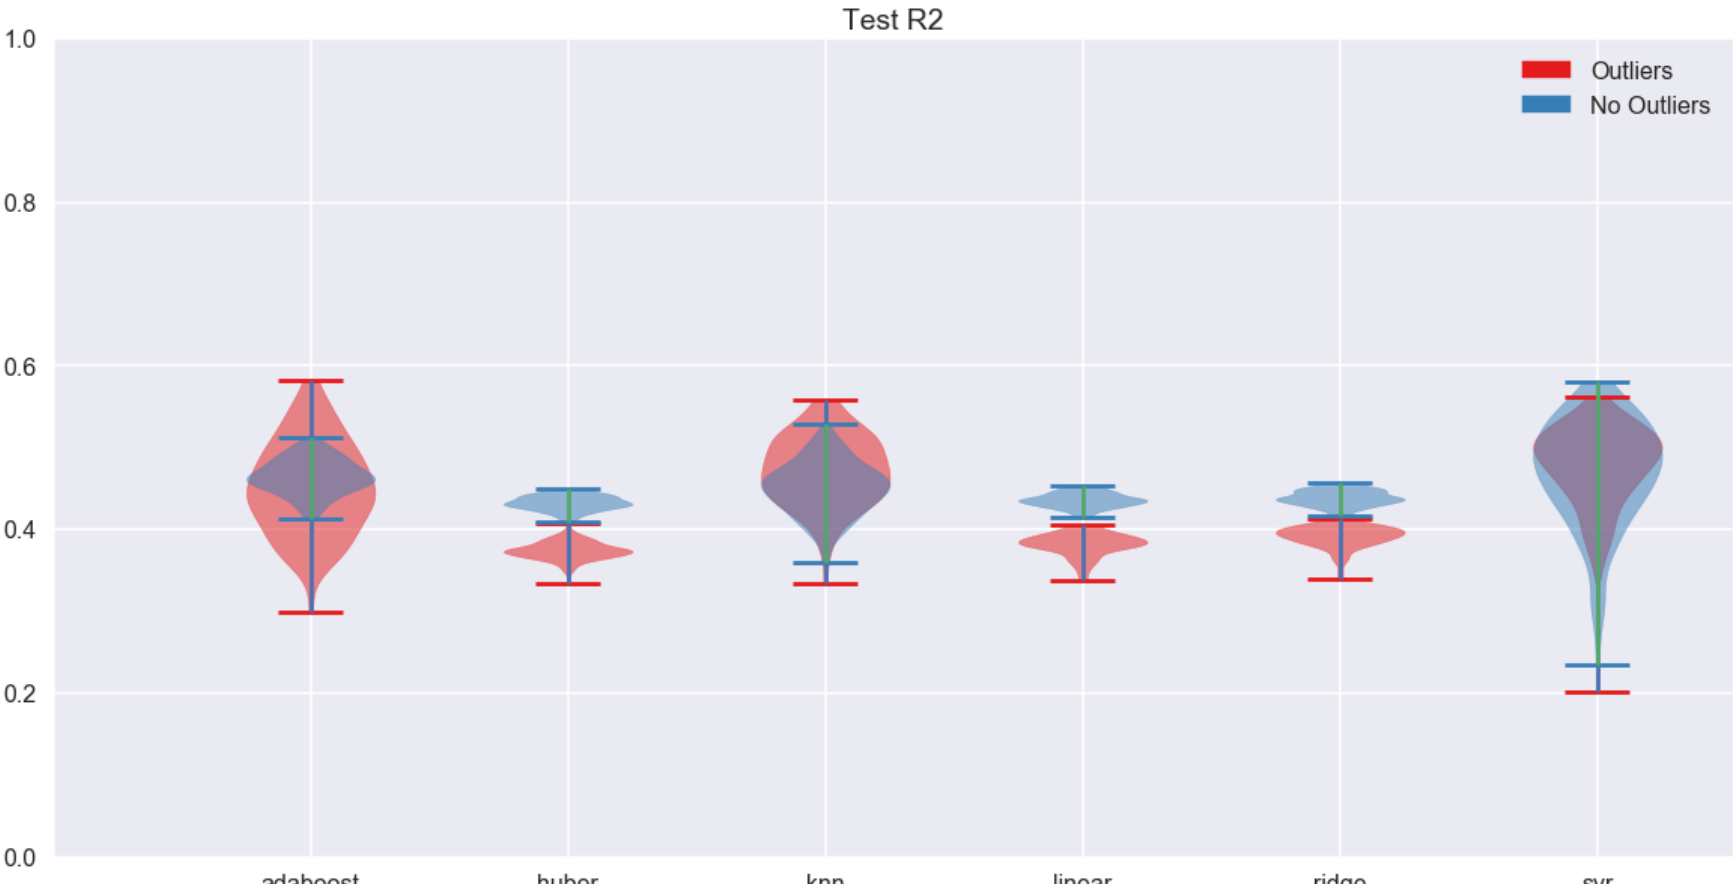
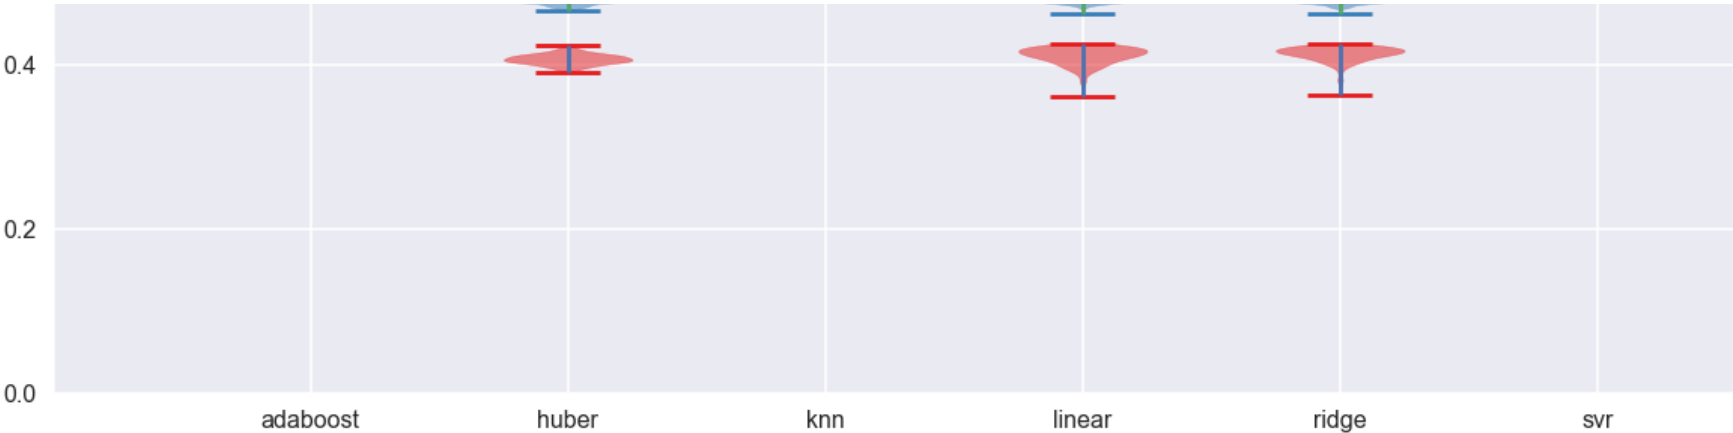
/Users/davidloving/anaconda3/envs/tf-gpu/lib/python3.6/site-packages/scipy/linalg/basic.py:1226:
RuntimeWarning: internal gelsd driver lwork query error, required iwork dimension not returned. This
is likely the result of LAPACK bug 0038, fixed in LAPACK 3.2.2 (released July 21, 2010). Falling
back to 'gelss' driver.
  warnings.warn(msg, RuntimeWarning)
```

Accuracy Results

```
In [11]: colors = plt.cm.Set1.colors

violin_plots([exp_1, exp_2],
             ['Train R2', 'Test R2'],
             experiment_name=['Outliers', 'No Outliers'],
             center_zero=False,
             cmap=colors)
```





Outlier Effects on Confidence Intervals

```
In [12]: # instantiate and fit models
def make_linear_models(x_train, y_train):
    md = dict()

    md['linear'] = LinearRegression().fit(x_train, y_train)
    md['ridge'] = RidgeCV(cv=15).fit(x_train, y_train)
    md['lasso'] = LassoCV(cv=15).fit(x_train, y_train)
    md['huber'] = GridSearchCV(HuberRegressor(), {'epsilon': [1.0, 1.2, 1.4, 1.6, 1.8]}).fit(x_train, y_train)

    return md
```

```
In [13]: coeff_names = df.drop('murder_per_100_k', axis=1).columns

# with outliers
exp_3 = run_experiment(make_linear_models, 100, x_train_all, y_train_all, x_test_all, y_test_all, coeff_names)

# without outliers
exp_4 = run_experiment(make_linear_models, 100, x_train_no, y_train_no, x_test_no, y_test_no, coeff_names)
```

Confidence Results

```
In [15]: violin_plots([exp_3, exp_4], coeff_names, experiment_name=['Outliers', 'No Outliers'], cmap=colors)
```



